

# Hackathon Setup guide



Leverage [featherless.ai](https://featherless.ai) - a serverless AI inference platform with unlimited access to thousands of open-source AI models. Focus on building amazing applications while we handle the infrastructure.

# Get Started in 3 Steps



## Sign up for Feather Premium

- Scan the QR code or click [here](#)
- **VSHACKS26** will automatically be entered
- Enjoy one month for free - you may need to refresh the page to see the changes

Scan here



Have fun!

### Feather Premium

Active

Renews on Mar 17, 2026

\$ 25 USD / Month Hackathon premium (100% off)

-  Access to DeepSeek, Kimi-K2 and GLM 4.6
-  Access any model - no limit on size!
-  Up to 4 concurrent connections
-  Up to 32K context

© 2026 Featherless.ai

 [Discord](#) [F.A.Q.](#) [About](#) [Terms](#) [Privacy](#)

[Cookies](#)

# Get Started in 3 Steps



## Get your API key

- After subscribing, go to the top right, click on your icon, and click on API Keys
- Here you can create an API key, keep this API key handy, you'll need it to access the service

The screenshot shows the Featherless web interface. The top navigation bar includes the Featherless logo, links for Models, Resources, Pricing, Chat, and Status, a search bar for models, and a user profile icon. The left sidebar contains a menu with sections for Featherless (API Keys, Private Models, Private Cloud) and Account (Profile, Change Password, Subscription, Payment Methods, Invoices, Referral Program). The main content area is titled 'API Keys' and includes a 'DEVELOPER ACCESS' badge. It features a 'New API Key' button and a table of active keys. The table has columns for KEY, CREATED, and LAST USED. Two keys are listed: a 'Secret Key' created on Jan 17, 2026, and an 'Agent: OpenClaw' key created on Mar 15, 2026. Both keys are partially masked with asterisks.

| KEY                                  | CREATED      | LAST USED    |
|--------------------------------------|--------------|--------------|
| Secret Key<br>rc_a0b34*****8d9d      | Jan 17, 2026 | Mar 20, 2026 |
| Agent: OpenClaw<br>rc_f4842*****97bc | Mar 15, 2026 | Mar 17, 2026 |



→ Test with a simple chat completion call!

Models

Resources

Pricing

Chat

Status

Search models...

Main

Architecture

Language

Domain

Training

License

Parameter Count

Reset

Filter parameter count...

7-9B 11.8k

1B 4.6k

<1B 3.3k

10-15B 3.2k

2-3B 2.6k

4-5B 2.1k

28-40B 831

+5 more

Model Family

Reset

Filter model family...

Llama 2 6.2k

Qwen 3 4.7k

Qwen 2 4.7k

Llama 3.2 3.1k

Mistral 2.6k

Llama 3.1 2.0k

Qwen 2.5 1.8k

+37 more

Capability

Reset

Filter capability...

Chat 440

Code 433

Models

29,775

Search models...

Popularity: High to Low

<< Previous

1

2

3

4

5

...

1489

Next >>

8B

8K

LLAMA3-8B

WARM

meta-llama/Meta-Llama-3-8B

6k · 3.4M · Apr 2024

8B

32K

LLAMA31-8B

WARM

meta-llama/Meta-Llama-3.1-8B-Instruct

6k · 8.3M

8B

32K

LLAMA31-8B

WARM

meta-llama/Llama-3.1-8B-Instruct

6k · 8.3M · Jul 2024

8B

8K

LLAMA3-8B

WARM

meta-llama/Meta-Llama-3-8B-Instruct

6k · 8.3M · Jul 2024

7B

4K

MISTRAL-V01-7B

WARM

mistralai/Mistral-7B-v0.1

6k · 8.3M · Jul 2024

3B

2K

PHI2-3B

COLD

microsoft/phi-2

6k · 8.3M · Jul 2024

# Popular Model Options



DeepSeek-V3.2

Advanced reasoning and coding capabilities



MiniMax-M2.5

Excellent in agentic tool use



Kimi-K2.5

Multimodal from the ground up



Mistral-Nemo-Instruct

Fast and efficient processing



GLM-5.1

Excels in long horizon tasks

# Your First API Call


## Using Python Requests

```
import requests
response = requests.post(

url="https://api.featherless.ai/v1/chat/completi
ons",
headers={
    "Authorization": "Bearer YOUR_API_KEY"
},
json={
    "model": "deepseek-ai/DeepSeek-V3-0324",
    "messages": [
        {"role": "user", "content": "Hello!"}
    ]
}
)
print(response.json())
```

## Using OpenAI SDK

```
from openai import OpenAI
client = OpenAI(
    base_url="https://api.featherless.ai/v1",
    api_key="YOUR_API_KEY"
)
response = client.chat.completions.create(
    model="deepseek-ai/DeepSeek-V3-0324",
    messages=[
        {"role": "user", "content": "Hello!"}
    ]
)
print(response.choices[0].message.content)
```



API

Data:

# Core API Endpoints

`/v1/chat/completions`

**Best for:** Chatbots, virtual assistants, conversational applications

- Structured messages with roles (system, user, assistant)
- Maintains conversation context automatically
- Ideal for interactive user-assistant interactions

`/v1/completions`

**Best for:** Content generation, text transformation, data extraction

- Takes a single prompt string
- Direct control over prompt format
- Maximum flexibility for custom use cases

# Troubleshooting Common Errors



## 401 - Unauthenticated

API key not recognized. Verify you've copied it correctly or generate a new one from account settings.



## 403 - Unauthorized

Model is gated. Visit the model's page, click "Unlock Model," and agree to license terms.



## 500 - Internal Server Error

Request could not be processed. Check for unsupported parameters in API documentation.



## 503 - Service Unavailable

Insufficient capacity or cold model. Retry the request. If it persists after three attempts, report on Discord.



# Concurrency Limits

Featherless operates on a unique model of **capacity reservation** rather than token-based billing. Your subscription tier determines the maximum size and number of concurrent AI model inference calls you can make, ensuring consistent performance without unpredictable token costs.

## Model Concurrency Costs

|                               |   |                               |
|-------------------------------|---|-------------------------------|
| 7B to 15B                     | 1 | Qwen 2.5 7B, Llama2 13B       |
| 24B to 34B                    | 2 | Qwen 32B Coder, Mistral 3 24B |
| 70B and 72B                   | 4 | Llama 3.3 70B, Qwen 2.5 72B   |
| Deepseek v3.2, R1 & Kimi-K2.5 | 4 | (Feather Premium only)        |

For a more advanced explanation and detailed examples, visit our [Concurrency Limits documentation](#).

# Application Guides

Explore our detailed application guides designed to help you seamlessly integrate Featherless.ai into your existing projects and workflows.

Find these guides and more on our documentation page:

[featherless.ai/docs/application-guides](https://featherless.ai/docs/application-guides)

## Coding

### Aider

Powering your in-terminal coding assistance with inference from featherless

### Cursor

Setting Up Cursor with Featherless.ai

### Roo Code

Roo Code (prev. Roo Cline) gives you a whole dev team of AI agents in your code editor.

### Cline

A fully collaborative AI partner that's open source, fully extensible, and designed to amplify developer impact.

## Workflow Automation

### n8n

Supercharge your AI workflows by connecting Featherless with n8n's powerful automation platform. Build sophisticated AI-powered automations that connect with 1000+ apps and services.

### Dify

Unleash 10,000+ open-source models in Dify with our dedicated plugin. Build powerful AI applications using serverless inference and low-code orchestration without the infrastructure headaches.

# Video walkthrough

Below is a video walkthrough that covers this PDF and additionally showcases a live use case of Featherless!



**Featherless AI – Quick Start and Live Demo.mov**



# Ready to Build!

You now have everything you need to start building amazing AI-powered applications. Here's your action plan:

## 1 Test Your Setup

Get your API key and test with a simple example

## 2 Explore Models

Find the right model for your use case in the catalog

## 3 Review Examples

Check the GitHub cookbook for code samples

## 4 Start Building

Begin your hackathon project with confidence

---

Quick Links: [Sign Up](#) · [Model Catalog](#) · [Documentation](#) · [GitHub](#)

**Good luck with your project!** We can't wait to see what you build!